

In [281...

```
%time
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from scipy.stats import zscore
```

CPU times: user 4 µs, sys: 0 ns, total: 4 µs
Wall time: 8.11 µs

## Data Cleaning

In [282...

```
df = pd.read_csv("../data/merged_data_cleaned.csv")
df.head()
```

Out[282...

	Unnamed: 0	Species	Owner	Country.of.Origin	Farm.Name	Lot.Number	Mi
0	0	Arabica	metad plc	Ethiopia	metad plc	NaN	meta p
1	1	Arabica	metad plc	Ethiopia	metad plc	NaN	meta p
2	2	Arabica	grounds for health admin	Guatemala	san marcos barrancas "san cristobal cuch	NaN	Na
3	3	Arabica	yidnekachew dabessa	Ethiopia	yidnekachew dabessa coffee plantation	NaN	wolens
4	4	Arabica	metad plc	Ethiopia	metad plc	NaN	meta p

5 rows × 44 columns

We have a lot of columns with data that are irrelevant for our analysis. We'll drop them to reduce dimensionality of the dataset

In [283...

```
df = df.drop(["Unnamed: 0", "Farm.Name", "Lot.Number", "Mill", "ICO.Numbe
```

In [284...

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1339 entries, 0 to 1338
Data columns (total 21 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Species                               1339 non-null   object
1   Country.of.Origin                     1338 non-null   object
2   Variety                               1113 non-null   object
3   Processing.Method                     1169 non-null   object
4   Aroma                                 1339 non-null   float64
5   Flavor                                1339 non-null   float64
6   Aftertaste                            1339 non-null   float64
7   Acidity                               1339 non-null   float64
8   Body                                  1339 non-null   float64
9   Balance                               1339 non-null   float64
10  Uniformity                            1339 non-null   float64
11  Clean.Cup                             1339 non-null   float64
12  Sweetness                             1339 non-null   float64
13  Cupper.Points                         1339 non-null   float64
14  Total.Cup.Points                     1339 non-null   float64
15  Moisture                             1339 non-null   float64
16  Category.One.Defects                 1339 non-null   int64
17  Quakers                              1338 non-null   float64
18  Color                                1069 non-null   object
19  Category.Two.Defects                 1339 non-null   int64
20  altitude_mean_meters                 1109 non-null   float64
dtypes: float64(14), int64(2), object(5)
memory usage: 219.8+ KB
```

```
In [285... df.isna().sum()
```

```
Out[285... Species                0
Country.of.Origin              1
Variety                        226
Processing.Method              170
Aroma                          0
Flavor                         0
Aftertaste                     0
Acidity                        0
Body                           0
Balance                        0
Uniformity                     0
Clean.Cup                      0
Sweetness                      0
Cupper.Points                  0
Total.Cup.Points               0
Moisture                       0
Category.One.Defects           0
Quakers                        1
Color                          270
Category.Two.Defects           0
altitude_mean_meters           230
dtype: int64
```

We replace rows with a lot of missing nominal data with the mode of the column, to retain the variance of the rest of the row data

```
In [286... df = df.fillna({'Variety': df['Variety'].mode()[0]})
df = df.fillna({'Processing.Method': df['Processing.Method'].mode()[0]})
```

```
df = df.fillna({'Color': df['Color'].mode()[0]})
```

We drop the rows with a single row missing nominal data. Especially "Country of Origin", since we can not just put in a mode value, since it might create significant wrong data

```
In [287... df = df.dropna(how='any')
```

```
In [288... df.isna().sum()
```

```
Out[288... Species                0
Country.of.Origin              0
Variety                       0
Processing.Method              0
Aroma                         0
Flavor                        0
Aftertaste                    0
Acidity                       0
Body                          0
Balance                       0
Uniformity                    0
Clean.Cup                     0
Sweetness                     0
Cupper.Points                 0
Total.Cup.Points              0
Moisture                      0
Category.One.Defects          0
Quakers                       0
Color                         0
Category.Two.Defects          0
altitude_mean_meters          0
dtype: int64
```

No more missing values!

## Outliers

```
In [289... df.describe()
```

```
Out[289... 
```

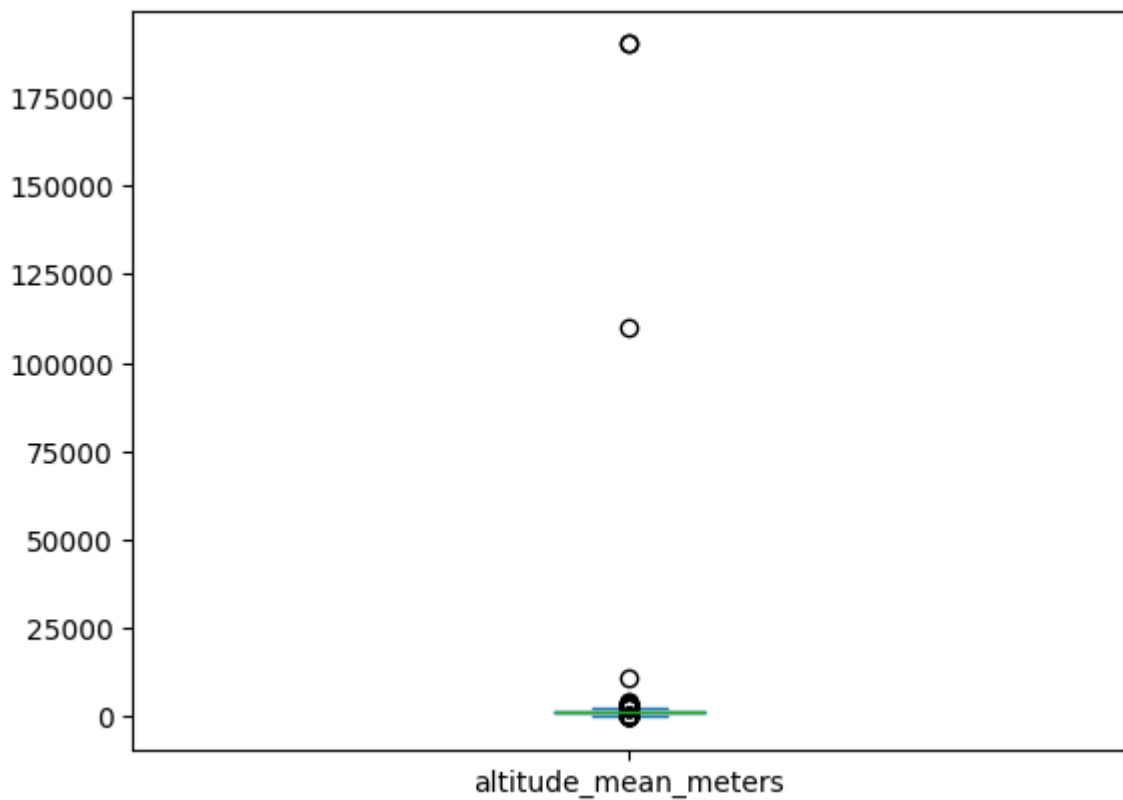
	Aroma	Flavor	Aftertaste	Acidity	Body	Balance
count	1108.000000	1108.000000	1108.000000	1108.000000	1108.000000	1108.000000
mean	7.570569	7.52056	7.394269	7.528953	7.506670	7.505542
std	0.383837	0.40059	0.405867	0.386075	0.366717	0.419311
min	0.000000	0.00000	0.000000	0.000000	0.000000	0.000000
25%	7.420000	7.33000	7.250000	7.330000	7.330000	7.330000
50%	7.580000	7.58000	7.420000	7.500000	7.500000	7.500000
75%	7.750000	7.75000	7.580000	7.750000	7.670000	7.750000
max	8.750000	8.83000	8.670000	8.750000	8.580000	8.750000

We have some columns with very high standard deviations

```
In [290... fig = plt.figure()
```

```
df.altitude_mean_meters.plot.box()
```

Out[290... <Axes: >



The column contains significant outliers. Since its only a couple of rows, we'll drop them

```
In [291... # We'll remove the outlying rows based on z-score  
df = df[np.abs(zscore(df['altitude_mean_meters'])) < 1]
```

```
In [292... df
```

Out[292...

	Species	Country.of.Origin	Variety	Processing.Method	Aroma	Flavor	Aftertaste
0	Arabica	Ethiopia	Caturra	Washed / Wet	8.67	8.83	8.67
1	Arabica	Ethiopia	Other	Washed / Wet	8.75	8.67	8.50
2	Arabica	Guatemala	Bourbon	Washed / Wet	8.42	8.50	8.42
3	Arabica	Ethiopia	Caturra	Natural / Dry	8.17	8.58	8.42
4	Arabica	Ethiopia	Other	Washed / Wet	8.25	8.50	8.25
...	...	...	...	...	...	...	...
1331	Robusta	India	Caturra	Washed / Wet	7.67	7.67	7.50
1332	Robusta	India	Caturra	Natural / Dry	7.58	7.42	7.42
1333	Robusta	United States	Arusha	Natural / Dry	7.92	7.50	7.42
1335	Robusta	Ecuador	Caturra	Washed / Wet	7.50	7.67	7.75
1336	Robusta	United States	Caturra	Natural / Dry	7.33	7.33	7.17

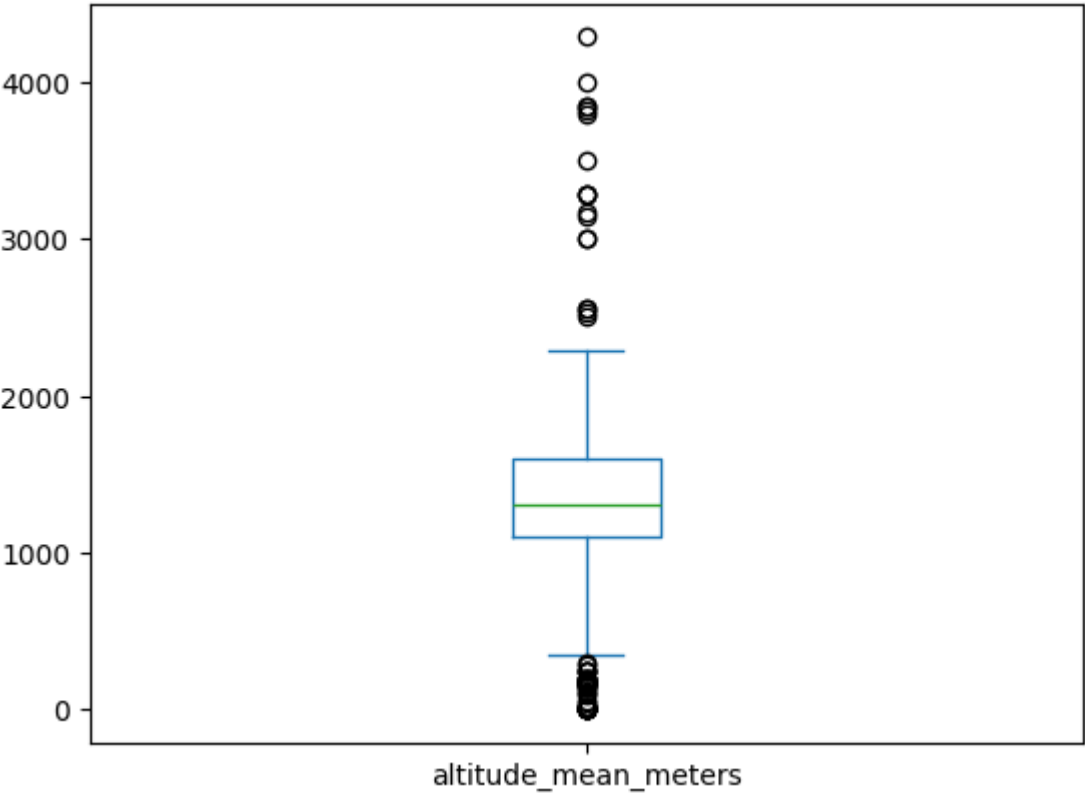
1104 rows × 21 columns

In [293...

df.altitude\_mean\_meters.plot.box()

Out[293...

<Axes: >




In [294...

df.plot.box(rot=90)

Out[294...


<Axes: >



Even after removing the worst outliers, the mean altitude still distributes over large values. We'll exclude it in the plot to identify other problematic features

```
In [295... df[df.columns.difference(['altitude_mean_meters'])].plot.box(rot=90)
```

```
Out[295... <Axes: >
```



It seems a lot of the features contain unnatural zero-values. We'll replace the zero values, with the median of the feature


```
In [296... df['Acidity'] = df['Acidity'].replace(0, df['Acidity'].median())
df['Aftertaste'] = df['Aftertaste'].replace(0, df['Aftertaste'].median())
df['Aroma'] = df['Aroma'].replace(0, df['Aroma'].median())
df['Balance'] = df['Balance'].replace(0, df['Balance'].median())
df['Body'] = df['Body'].replace(0, df['Body'].median())
df['Clean.Cup'] = df['Clean.Cup'].replace(0, df['Clean.Cup'].median())
df['Cupper.Points'] = df['Cupper.Points'].replace(0, df['Cupper.Points'].median())
df['Flavor'] = df['Flavor'].replace(0, df['Flavor'].median())
df['Moisture'] = df['Moisture'].replace(0, df['Moisture'].median())
df['Sweetness'] = df['Sweetness'].replace(0, df['Sweetness'].median())
df['Uniformity'] = df['Uniformity'].replace(0, df['Uniformity'].median())
```

Except the Total cup points. We'll drop the row since it is our target value, and an unnatural zero might mess with correlations

```
In [297... df = df[df['Total.Cup.Points'] != 0]
```

```
In [298... df[df.columns.difference(['altitude_mean_meters'])].plot.box(rot=90)
```

```
Out[298... <Axes: >
```



```
In [299... df.describe()
```

```
Out[299...

```

	Aroma	Flavor	Aftertaste	Acidity	Body	Balance
count	1103.000000	1103.000000	1103.000000	1103.000000	1103.000000	1103.000000
mean	7.578368	7.527824	7.401496	7.535739	7.513654	7.512660
std	0.309181	0.331268	0.340065	0.313192	0.289467	0.354020
min	5.080000	6.170000	6.170000	5.250000	5.170000	5.250000
25%	7.420000	7.330000	7.250000	7.330000	7.330000	7.330000
50%	7.580000	7.580000	7.420000	7.500000	7.500000	7.500000
75%	7.750000	7.750000	7.580000	7.750000	7.670000	7.750000
max	8.750000	8.830000	8.670000	8.750000	8.580000	8.750000

```
In [300... df.to_csv('cleaned_dataset_no_zeros.csv', index=False)
```

The dataset is now ready for analysis!